

lab07_music_partial

April 5, 2018

1 Lab 7: Neural Networks for Music Classification

In addition to the concepts in the [MNIST neural network demo](#), in this lab, you will learn to:

- * Load a file from a URL
- * Extract simple features from audio samples for machine learning tasks such as speech recognition and classification
- * Build a simple neural network for music classification using these features
- * Use a callback to store the loss and accuracy history in the training process
- * Optimize the learning rate of the neural network

To illustrate the basic concepts, we will look at a relatively simple music classification problem. Given a sample of music, we want to determine which instrument (e.g. trumpet, violin, piano) is playing. This dataset was generously supplied by [Prof. Juan Bello](#) at NYU Stenhardt and his former PhD student Eric Humphrey (now at Spotify). They have a complete website dedicated to deep learning methods in music informatics:

<http://marl.smusic.nyu.edu/wordpress/projects/feature-learning-deep-architectures/deep-learning-python-tutorial/>

You can also check out Juan's course.

1.1 Loading the Keras package

We begin by loading keras and the other packages

```
In [1]: import keras
```

```
/usr/local/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the s
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
In [2]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

1.2 Audio Feature Extraction with Librosa

The key to audio classification is to extract the correct features. In addition to keras, we will need the librosa package. The librosa package in python has a rich set of methods extracting the features of audio samples commonly used in machine learning tasks such as speech recognition and sound classification.

Installation instructions and complete documentation for the package are given on the [librosa main page](#). On most systems, you should be able to simply use:

```
pip install -u librosa
```

For Unix, you may need to load some additional packages:

```
sudo apt-get install build-essential
sudo apt-get install libxext-dev python-qt4 qt4-dev-tools
pip install librosa
```

After you have installed the package, try to import it.

```
In [3]: import librosa
import librosa.display
import librosa.feature
```

In this lab, we will use a set of music samples from the website:

<http://theremin.music.uiowa.edu>

This website has a great set of samples for audio processing. Look on the web for how to use the `requests.get` and `file.write` commands to load the file at the URL provided into your working directory.

You can play the audio sample by copying the file to your local machine and playing it on any media player. If you listen to it you will hear a soprano saxophone (with vibrato) playing four notes (C, C#, D, Eb).

```
In [4]: import requests
fn = "SopSax.Vib.pp.C6Eb6.aiff"
url = "http://theremin.music.uiowa.edu/sound files/MIS/Woodwinds/sopranosaxophone/"+fn

# TODO: Load the file from url and save it in a file under the name fn
r = requests.get(url)
wf = open(fn, 'wb')
wf.write(r.content)
wf.close()
```

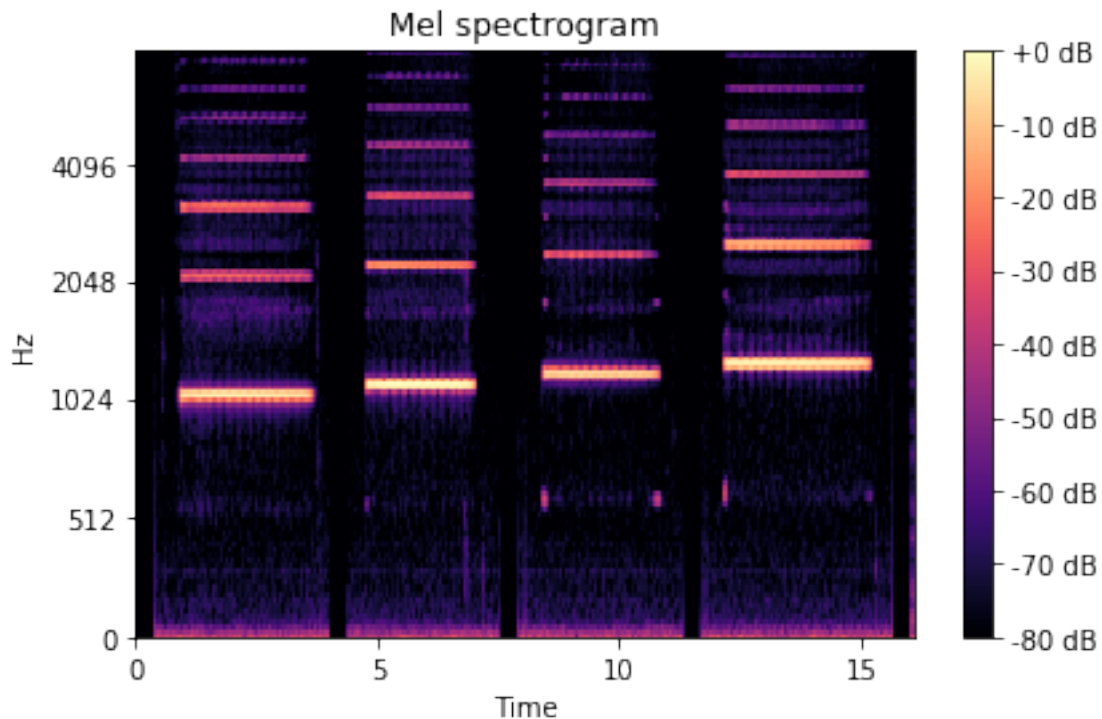
Next, use librosa command `librosa.load` to read the audio file with filename `fn` and get the samples `y` and sample rate `sr`.

```
In [5]: # TODO
y, sr = librosa.load(fn)
```

Extracting features from audio files is an entire subject on its own right. A commonly used set of features are called the Mel Frequency Cepstral Coefficients (MFCCs). These are derived from the so-called mel spectrogram which is something like a regular spectrogram, but the power and frequency are represented in log scale, which more naturally aligns with human perceptual processing. You can run the code below to display the mel spectrogram from the audio sample.

You can easily see the four notes played in the audio track. You also see the ‘harmonics’ of each notes, which are other tones at integer multiples of the fundamental frequency of each note.

```
In [6]: S = librosa.feature.melspectrogram(y=y, sr=sr, n_mels=128, fmax=8000)
librosa.display.specshow(librosa.power_to_db(S,ref=np.max),
                        y_axis='mel', fmax=8000, x_axis='time')
plt.colorbar(format='%+2.0f dB')
plt.title('Mel spectrogram')
plt.tight_layout()
```



1.3 Downloading the Data

Using the MFCC features described above, Eric Humphrey and Juan Bellow have created a complete data set that can be used for instrument classification. Essentially, they collected a number of data files from the website above. For each audio file, they segmented the track into notes and then extracted 120 MFCCs for each note. The goal is to recognize the instrument from the 120 MFCCs. The process of feature extraction is quite involved. So, we will just use their processed data provided at:

<https://github.com/marl/dl4mir-tutorial/blob/master/README.md>

Note the password. Load the four files into some directory, say `instrument_dataset`. Then, load them with the commands.

```
In [7]: data_dir = 'instrument_dataset/'
Xtr = np.load(data_dir+'uiowa_train_data.npy')
ytr = np.load(data_dir+'uiowa_train_labels.npy')
Xts = np.load(data_dir+'uiowa_test_data.npy')
yts = np.load(data_dir+'uiowa_test_labels.npy')
```

Looking at the data files: * What are the number of training and test samples? * What is the number of features for each sample? * How many classes (i.e. instruments) are there per class.

```
In [8]: # TODO
print("Number of training samples: %d" % ytr.shape[0])
print("Number of test samples: %d" % yts.shape[0])
print("Number of features: %d" % Xtr.shape[1])
print("Classes: %d" % np.unique(yts).shape[0])
```

```
Number of training samples: 66247
Number of test samples: 14904
Number of features: 120
Classes: 10
```

Before continuing, you must scale the training and test data, `Xtr` and `Xts`. Compute the mean and std deviation of each feature in `Xtr` and create a new training data set, `Xtr_scale`, by subtracting the mean and dividing by the std deviation. Also compute a scaled test data set, `Xts_scale` using the mean and std deviation learned from the training data set.

```
In [9]: # TODO Scale the training and test matrices
Xtr_scale = (Xtr - Xtr.mean(axis=0)) / Xtr.std(axis=0)
Xts_scale = (Xts - Xts.mean(axis=0)) / Xts.std(axis=0)
```

1.4 Building a Neural Network Classifier

Following the example in [MNIST neural network demo](#), clear the keras session. Then, create a neural network model with: * `nh=256` hidden units * sigmoid activation for the hidden layer, softmax for the output layer * select the input and output shapes correctly * print the model summary

```
In [10]: from keras.models import Model, Sequential
         from keras.layers import Dense, Activation

In [11]: # TODO clear session
         import keras.backend as K
         K.clear_session()

In [12]: # TODO: construct the model
         nh = 256
         nin = Xtr_scale.shape[1]
         nout = np.unique(yts).shape[0]

         model = Sequential()
         model.add(Dense(nh, input_shape=(nin,), activation='sigmoid'))
         model.add(Dense(nout, activation='softmax'))

In [13]: # TODO: Print the model summary
         model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	30976
dense_2 (Dense)	(None, 10)	2570
Total params: 33,546		
Trainable params: 33,546		
Non-trainable params: 0		

To keep track of the loss history and validation accuracy, we can use a *callback* function as described in [Keras callback documentation](#). A callback is a class that is passed to the `fit` method. Here we provide the definition of `LossHistory` callback class below to save the loss and validation accuracy. This callback class allows you to record the loss at the batch level in addition to at the epoch level. However, for this lab, you could choose to just use the returned history class by `model.fit`, which will allow you to plot the metrics at the epoch level. For your own practice, you could use this callback class instead of the returned history class to plot the results required below.

```
In [14]: class LossHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        # TODO: Create four empty lists, self.loss, self.acc, self.val_acc and self
        self.loss = []
        self.acc = []
        self.val_acc = []
        self.batch_loss = []

    def on_batch_end(self, batch, logs={}):
        # TODO: This is called at the end of each batch.
        # Add the loss in logs.get('loss') to the batch_loss list
        self.batch_loss.append(logs.get('loss'))

    def on_epoch_end(self, epoch, logs):
        # TODO: This is called at the end of each epoch.
        # Add the training accuracy in logs.get('acc') to the acc list
        # Add the test accuracy in logs.get('val_acc') to the val_acc list
        # Add the training loss in logs.get('loss') to the loss list

        self.acc.append(logs.get('acc'))
        self.val_acc.append(logs.get('val_acc'))
        self.loss.append(logs.get('loss'))

    # Create an instance of the history callback
    history_cb = LossHistory()
```

Create an optimizer and compile the model. Select the appropriate loss function and metrics. For the optimizer, use the Adam optimizer with a learning rate of 0.001

```
In [15]: # TODO
        from keras import optimizers
        opt = optimizers.Adam(lr=0.001)
        model.compile(opt, loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Fit the model for 10 epochs using the scaled data for both the training and validation. Use the `validation_data` option to pass the test data. Also, use the return from `model.fit` to record the history of loss, accuracy, and validation accuracy in successive epochs. Use a batch size of 100. Your final accuracy should be >99%. If you want, you could also use the callback class you defined to record the history.

```
In [16]: # TODO
        batch_size = 100
        epochs = 10
        model.fit(Xtr_scale, ytr, batch_size=batch_size, epochs=epochs, callbacks=[history_cb])
```

Train on 66247 samples, validate on 14904 samples

```
Epoch 1/10
66247/66247 [=====] - 1s 23us/step - loss: 0.3667 - acc: 0.8987 - val.
Epoch 2/10
66247/66247 [=====] - 1s 20us/step - loss: 0.1034 - acc: 0.9750 - val.
Epoch 3/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0604 - acc: 0.9854 - val.
Epoch 4/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0425 - acc: 0.9893 - val.
Epoch 5/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0324 - acc: 0.9914 - val.
Epoch 6/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0253 - acc: 0.9935 - val.
Epoch 7/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0213 - acc: 0.9945 - val.
Epoch 8/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0174 - acc: 0.9953 - val.
Epoch 9/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0149 - acc: 0.9962 - val.
Epoch 10/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0131 - acc: 0.9965 - val.
```

```
Out[16]: <keras.callbacks.History at 0x10810bc88>
```

Plot the training loss, accuracy and validation accuracy vs. epoch, using the returned history record from `model.fit`. You should produce two subplots. One subplot contains the curve of loss vs. epochs. The other subplot contains two curves, one for the training accuracy and another for the validation accuracy. You should see that the loss continuously decreases, the accuracy continuously increases, but the validation accuracy saturates at a little higher than 99%. After that it “bounces around” due to the noise in the stochastic gradient descent.

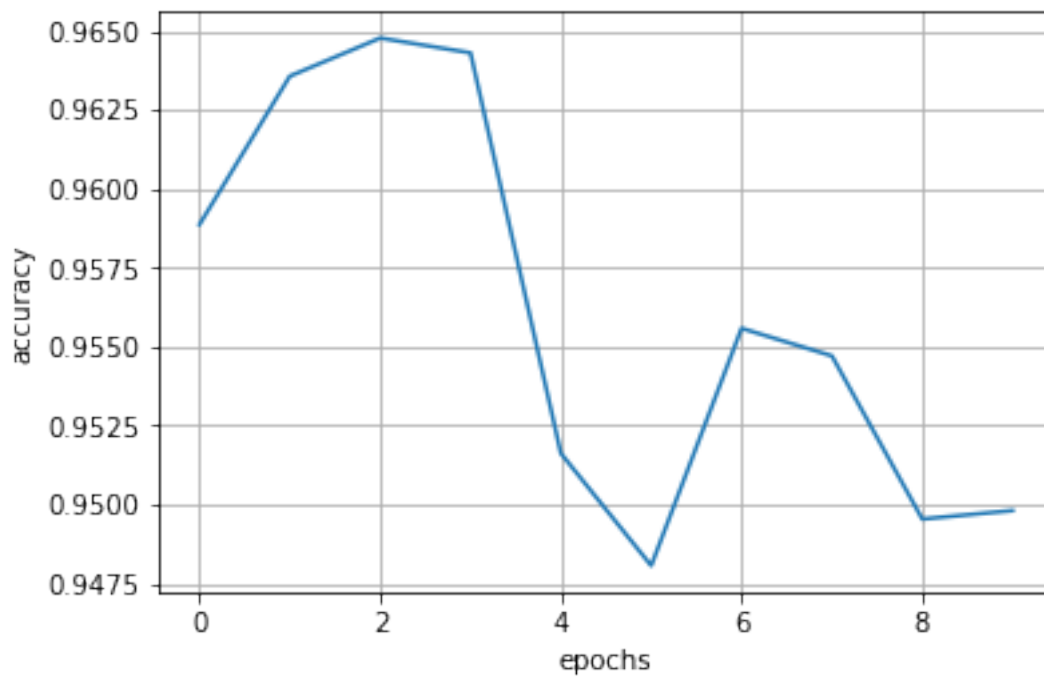
You could also try to plot the loss values saved in the `history_cb` class. In addition to plot the metrics for every epoch, you could plot the `batch_loss` per batch. But you may want to plot the

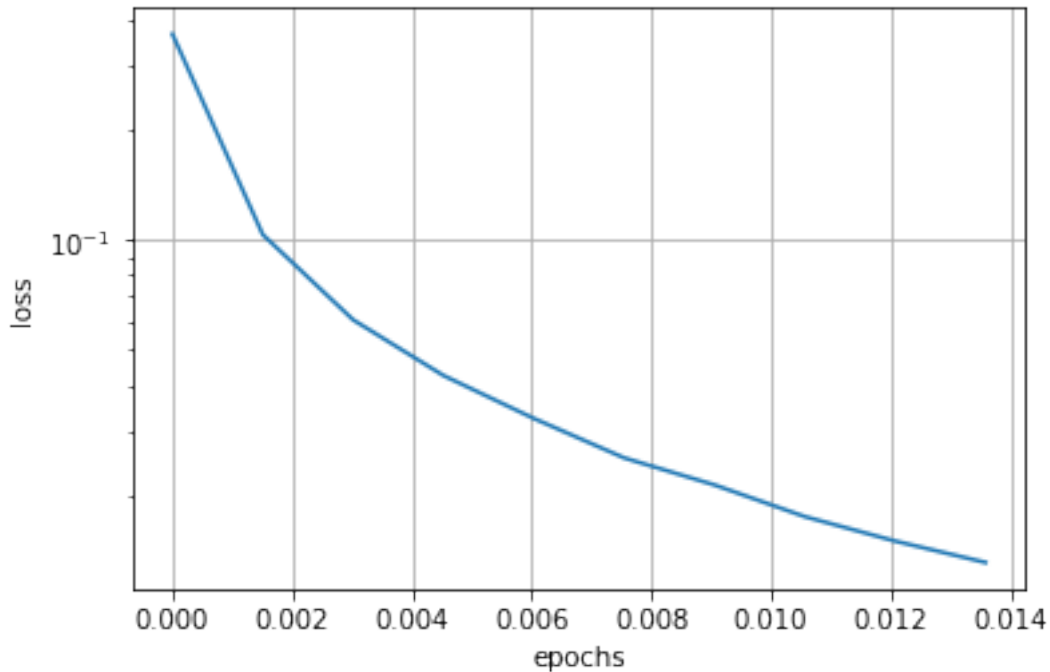
x-axis in epochs. Note that the epoch in step i is $\text{epoch} = i * \text{batch_size} / \text{ntr}$ where batch_size is the batch_size and ntr is the total number of training samples.

```
In [17]: # TODO
plt.plot(history_cb.val_acc)
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.grid()
plt.show()

step = np.arange(len(history_cb.loss))
step = step * batch_size / ytr.shape[0]

plt.semilogy(step, history_cb.loss)
plt.xlabel('epochs')
plt.ylabel('loss')
plt.grid()
```





1.5 Optimizing the Learning Rate

One challenge in training neural networks is the selection of the learning rate. Rerun the above code, trying three learning rates as shown in the vector rates. For each learning rate: * clear the session * construct the network * select the optimizer. Use the Adam optimizer with the appropriate learning rate. * train the model * save the accuracy and losses

```
In [18]: rates = [0.01,0.001,0.0001]
        batch_size = 100
        loss_hist = []
        val_acc_hist = []

        # TODO
        nh = 256
        nin = Xtr_scale.shape[1]
        nout = np.unique(yts).shape[0]

        # for each learning rate
        for rate in rates:

            # clearing the session
            K.clear_session()

            # constructing the network
            model = Sequential()
```



```

model.add(Dense(nh, input_dim=nin))
model.add(Activation('sigmoid'))
model.add(Dense(nout))
model.add(Activation('softmax'))

# selecting the optimizer
opt = optimizers.Adam(lr=rate)
model.compile(opt, loss='sparse_categorical_crossentropy', metrics=['accuracy'])

#train the model
model.fit(Xtr_scale, ytr, batch_size=batch_size, epochs=10, callbacks=[history_cb])

# saving accuracy and losses
loss_hist.append(history_cb.loss)
val_acc_hist.append(history_cb.val_acc)

```

Train on 66247 samples, validate on 14904 samples

```

Epoch 1/10
66247/66247 [=====] - 1s 22us/step - loss: 0.1018 - acc: 0.9686 - val.
Epoch 2/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0278 - acc: 0.9912 - val.
Epoch 3/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0205 - acc: 0.9933 - val.
Epoch 4/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0198 - acc: 0.9934 - val.
Epoch 5/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0161 - acc: 0.9949 - val.
Epoch 6/10
66247/66247 [=====] - 1s 19us/step - loss: 0.0194 - acc: 0.9944 - val.
Epoch 7/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0129 - acc: 0.9955 - val.
Epoch 8/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0111 - acc: 0.9966 - val.
Epoch 9/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0120 - acc: 0.9962 - val.
Epoch 10/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0144 - acc: 0.9956 - val.

```

Train on 66247 samples, validate on 14904 samples

```

Epoch 1/10
66247/66247 [=====] - 2s 23us/step - loss: 0.3606 - acc: 0.9026 - val.
Epoch 2/10
66247/66247 [=====] - 1s 20us/step - loss: 0.1017 - acc: 0.9751 - val.
Epoch 3/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0598 - acc: 0.9855 - val.
Epoch 4/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0422 - acc: 0.9893 - val.
Epoch 5/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0318 - acc: 0.9918 - val.

```

```

Epoch 6/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0254 - acc: 0.9933 - val.
Epoch 7/10
66247/66247 [=====] - 1s 21us/step - loss: 0.0203 - acc: 0.9944 - val.
Epoch 8/10
66247/66247 [=====] - 1s 21us/step - loss: 0.0171 - acc: 0.9958 - val.
Epoch 9/10
66247/66247 [=====] - 1s 21us/step - loss: 0.0148 - acc: 0.9963 - val.
Epoch 10/10
66247/66247 [=====] - 1s 20us/step - loss: 0.0130 - acc: 0.9967 - val.
Train on 66247 samples, validate on 14904 samples
Epoch 1/10
66247/66247 [=====] - 2s 25us/step - loss: 1.1038 - acc: 0.6649 - val.
Epoch 2/10
66247/66247 [=====] - 1s 21us/step - loss: 0.5412 - acc: 0.8548 - val.
Epoch 3/10
66247/66247 [=====] - 1s 22us/step - loss: 0.3732 - acc: 0.9136 - val.
Epoch 4/10
66247/66247 [=====] - 1s 22us/step - loss: 0.2885 - acc: 0.9354 - val.
Epoch 5/10
66247/66247 [=====] - 1s 21us/step - loss: 0.2346 - acc: 0.9478 - val.
Epoch 6/10
66247/66247 [=====] - 1s 21us/step - loss: 0.1960 - acc: 0.9554 - val.
Epoch 7/10
66247/66247 [=====] - 1s 20us/step - loss: 0.1666 - acc: 0.9612 - val.
Epoch 8/10
66247/66247 [=====] - 1s 22us/step - loss: 0.1433 - acc: 0.9662 - val.
Epoch 9/10
66247/66247 [=====] - 1s 20us/step - loss: 0.1246 - acc: 0.9705 - val.
Epoch 10/10
66247/66247 [=====] - 1s 21us/step - loss: 0.1095 - acc: 0.9742 - val.

```

Plot the loss function vs. the epoch number for all three learning rates on one graph. You should see that the lower learning rates are more stable, but converge slower. Similarly, plot the accuracy and validation accuracy and make your observations. You may want to plot the three set of figures as three subplots.

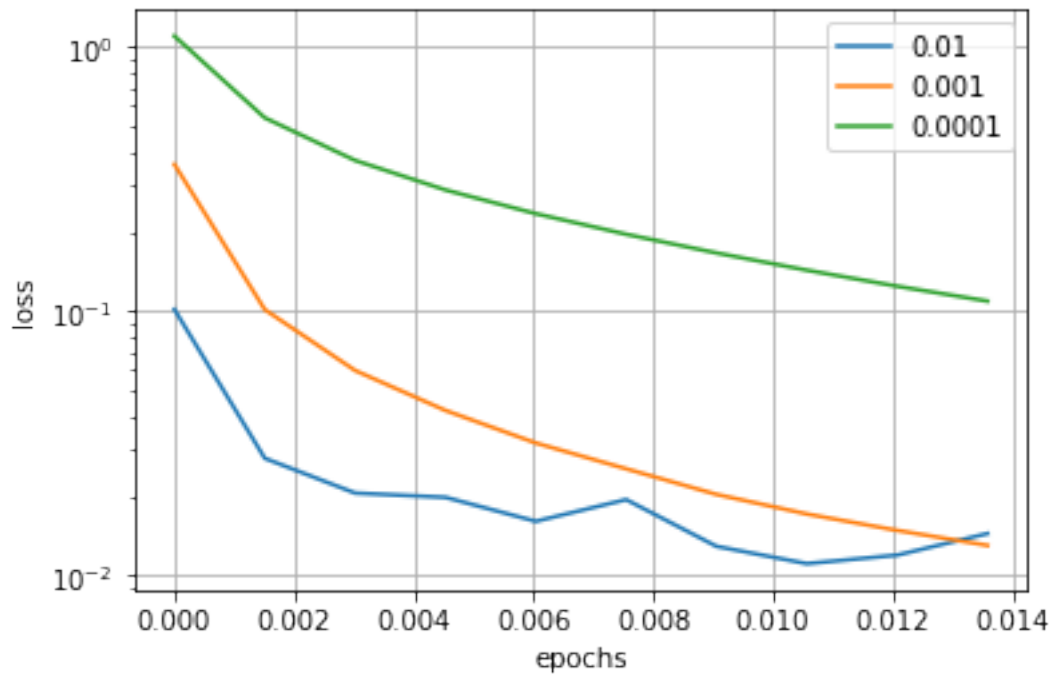
```

In [19]: # TODO
         step = np.arange(len(history_cb.loss))
         step = step * batch_size / ytr.shape[0]

         plt.grid()
         plt.xlabel('epochs')
         plt.ylabel('loss')
         plt.semilogy(step, loss_hist[0])
         plt.semilogy(step, loss_hist[1])
         plt.semilogy(step, loss_hist[2])
         plt.legend((rates[0], rates[1], rates[2]))

```

Out[19]: <matplotlib.legend.Legend at 0x11f033198>



Question: Which learning rate is the best ?

Answer: From the figure above, we can see that the loss of a lower learning rate is more stable, but it converges slower. So generally, a lower learning rate is better to reach a stable result. However, we could choose larger learning rate to help us to save time in training. For example, in this case, 0.001 may be the best learning rate, since it balanced the stability and the rate of convergence.